Assignment 4

Name – Sati, Ankit Date – 3/17/2022

Section – 001

SID – as14128

Total in points (Maximum 100 points)–

Professors Comments –

Affirmation of Independent Effort – Ankit Sati

Index

Note to Grader - We have discussed with the professor that we only have to complete **4 of the 7 topics** in the assignment.

However I have done 6 and have only missed part 4.

Have discussed the same with professor as **HDInsight is a paid service.**

Part I - Completed

Part II - Completed

Part V - Completed

Part III - Completed

Part VI - Completed

Part VII - Completed

Part VII - HDinsight is a paid service – left this one. (Explained above)

Part I – Get Familiar with spark.

1. **Setting up resources to run the Spark Euler.**

```
ankit@LAPTOP-S2U1QMGB:/mnt/c$ docker run -it -p 8888:8888 dbgannon/tutorial
cp: omitting directory '/tutorial_notebooks/graph'
/home/jovyan/work
total 16
drwxr-xr-x 1 jovyan users 4096 Mar 24 23:31 .
drwxr-xr-x 1 jovyan users 4096 Jul  8  2017 ..
drwxr-xr-x 2 jovyan users 4096 Mar 24 23:31 notebooks
[I 23:31:53.041 NotebookApp] Writing notebook server cookie secret to /home/jovyan/.local/share/jupyter/runtime/notebook_cookie
_secret
[I 23:31:53.714 NotebookApp] JupyterLab alpha preview extension loaded from /opt/conda/lib/python3.5/site-packages/jupyterlab
[I 23:31:53.720 NotebookApp] Serving notebooks from local directory: /home/jovyan/work
[I 23:31:53.720 NotebookApp] 0 active kernels
[I 23:31:53.721 NotebookApp] The Jupyter Notebook is running at: https://[all ip addresses on your system]:8888/
```

2. **Setting up the Eular-Spark in the same directory.**
   - **$ export PYSPARK_DRIVER_PYTHON=Jupyter**
   - **$ export PYSPARK_DRIVER_PYTHON_OPTS=notebook**

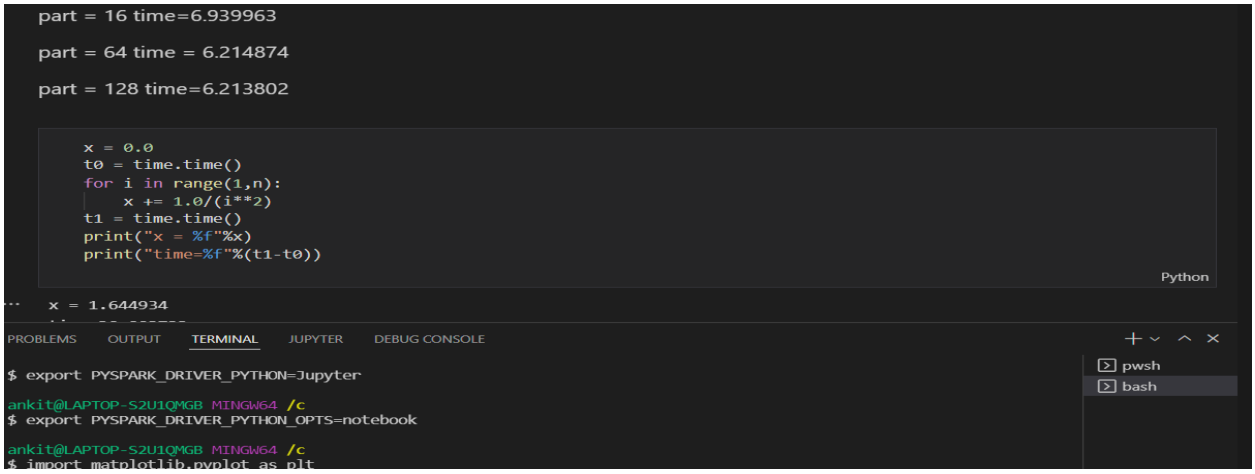

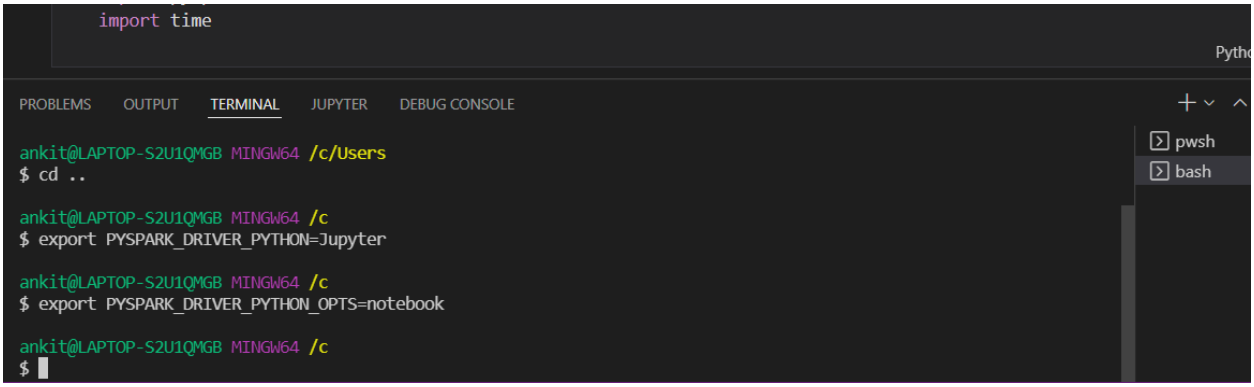

Second Notebook just to experiment with. I choose SQL-Magic

- pip install pandas
- pip install sqlalchemy # ORM for databases

Exampe 2. A K-means computation in Spark

This is a simple demo of using spark to compute k-means where k = 4. To make it easy to repeat the computation with different numbers of point and check the accuracy and performance, we create an artificial set of points in the plane where there are 4 clusters of "random" points.

we first import the python spark package and several others we will need. If pyspark does not load you may have the wrong kernel running. On the data science vm, look for kernel "Spark-python".

[1]:
```
import sys
import time
import pyspark
from random import random
```

we are going to do some plotting, so let's set that up to. also let's get the numpy library and call it np.

[2]:
```
import matplotlib.pyplot as plt
%matplotlib inline
```

```
Matplotlib is building the font cache; this may take a moment.
```

[3]:
```
import numpy as np
```

we first create the set of "n" points. Later we will come back and try different values of n, but this is a good place to start. We create n pairs of the form (0.0, 0.0)

[4]:
```
n = 100000
nums = np.zeros((n,2))
```

[5]:
```
nums.shape
```

[5]:
```
(100000, 2)
```

Now we create four "random" clusters. Each cluster is in the shape of a small circle of points. This is so that we can repeat the experiment and it will converge in the same way each time.

[6]:
```
for i in range(int(n/4)):
    x = random()*3.1415*2
    s =  0.6
    ranx = s*(2*random()-1)
    rany = s*(2*random()-1)
```

```
ankit@LAPTOP-S2U1QMGB:/mnt/c$ pip install sqlalchemy # ORM for databases
Collecting sqlalchemy
  Downloading SQLAlchemy-1.4.32-cp38-cp38-manylinux_2_5_x86_64.manylinux1_x86_64.manylinux_2_17_x86_64.manylinux
     |████████████████████████████████| 1.6 MB 1.1 MB/s
Collecting greenlet!=0.4.17; python_version >= "3" and (platform_machine == "aarch64" or (platform_machine == "p
"AMD64" or (platform_machine == "win32" or platform_machine == "WIN32"))))))
  Downloading greenlet-1.1.2-cp38-cp38-manylinux_2_17_x86_64.manylinux2014_x86_64.whl (156 kB)
     |████████████████████████████████| 156 kB 3.2 MB/s
Installing collected packages: greenlet, sqlalchemy
Successfully installed greenlet-1.1.2 sqlalchemy-1.4.32
```

- pip install ipython-sql # SQL magic function

```
ankit@LAPTOP-S2U1QMGB:/mnt/c$ pip install ipython-sql # SQL magic function
Collecting ipython-sql
  Downloading ipython_sql-0.4.0-py3-none-any.whl (19 kB)
Requirement already satisfied: sqlalchemy>=0.6.7 in /home/ankit/.local/lib/python3.8/site-packages (from ipyth
Requirement already satisfied: ipython>=1.0 in /home/ankit/.local/lib/python3.8/site-packages (from ipython-sq
Collecting prettytable<1
  Downloading prettytable-0.7.2.tar.bz2 (21 kB)
Requirement already satisfied: six in /usr/lib/python3/dist-packages (from ipython-sql) (1.14.0)
Requirement already satisfied: ipython-genutils>=0.1.0 in /home/ankit/.local/lib/python3.8/site-packages (from
Collecting sqlparse
  Downloading sqlparse-0.4.2-py3-none-any.whl (42 kB)
     |████████████████████████████████| 42 kB 967 kB/s
Requirement already satisfied: greenlet!=0.4.17; python_version >= "3" and (platform_machine == "aarch64" or (p
platform_machine == "AMD64" or (platform_machine == "win32" or platform_machine == "WIN32")))))) in /home/ankit
Requirement already satisfied: stack-data in /home/ankit/.local/lib/python3.8/site-packages (from ipython>=1.0-
```

```
    return sqlContext.sql(val).show(max_show_lines)

@register_line_cell_magic
def sql_display(line, cell=None):
    """Execute sql and convert results to Pandas DataFrame for pretty display or further processing.
    Use: %sql_display or %%sql_display"""
    val = cell if cell is not None else line
    return sqlContext.sql(val).limit(max_show_lines).toPandas()

@register_line_cell_magic
def sql_explain(line, cell=None):
    "Display the execution plan of the sql. Use: %sql_explain or %%sql_explain"
    val = cell if cell is not None else line
    return sqlContext.sql(val).explain(detailed_explain)
```

Python

```
sqlCtx = SQLContext(sc)
```

Python

```
import csv
with open("/Users/dennisgannon/Desktop/hvac.csv", 'rb') as csvfile:
    spamreader = csv.reader(csvfile) #, delimiter=',', quotechar='|')
    for row in spamreader:
        print ', '.join(row)
```

Python

```
3/23/2016, 11:45, 67, 54, headquarters
3/23/2016, 11:51, 67, 77, lab1
```

BLEMS OUTPUT TERMINAL JUPYTER DEBUG CONSOLE

pwsh
bash

3. **Sql-magic**

Finally now we can go ahead and setup the SQL magic over the Jupyter Notebook.
We will be using the same notebook as shared in the example in the tutorial.

File Edit View Insert Cell Kernel Help Trusted Python 3 (ipykernel)

Run Code

```
6/6/13, 17:43:51, 67, 79, 1, 9, 10
6/7/13, 18:43:51, 67, 72, 13, 29, 15
6/8/13, 19:43:51, 69, 76, 11, 25, 11
6/9/13, 20:43:51, 65, 67, 6, 21, 16
6/10/13, 21:43:51, 67, 74, 3, 19, 14
6/11/13, 22:43:51, 68, 76, 9, 28, 2
6/12/13, 23:43:51, 68, 69, 7, 26, 20
6/13/13, 0:13:20, 67, 77, 8, 19, 12
6/14/13, 1:13:20, 70, 68, 6, 1, 1
6/15/13, 2:13:20, 70, 57, 7, 12, 6
6/16/13, 3:13:20, 67, 55, 11, 16, 12
6/17/13, 4:13:20, 70, 55, 17, 17, 19
6/18/13, 5:13:20, 70, 58, 6, 17, 19
6/19/13, 6:13:20, 69, 75, 9, 21, 6
6/20/13, 7:13:20, 69, 75, 13, 14, 16
6/21/13, 8:13:20, 69, 55, 1, 8, 14
6/22/13, 9:13:20, 65, 68, 15, 20, 3
6/23/13, 10:13:20, 65, 63, 5, 22, 19
```

```
In [8]: hvacText = sc.textFile("/Users/hemantp/Download/hvac.csv")

In [9]: chema = StructType([StructField("date", StringType(), False),StructField("time", StringType(), False),StructField("targ

In [10]: hvacSchema

Out[10]: StructType(List(StructField(date,StringType,false),StructField(time,StringType,false),StructField(targettemp,IntegerT
         ype,false),StructField(actualtemp,IntegerType,false),StructField(buildingID,StringType,false)))

In [11]: hvac = hvacText.map(lambda s: s.split(",")).filter(lambda s: s[0] != "Date").map(lambda s:(str(s[0]), str(s[1]), int(s[

In [12]: hvac.collect()
```

```
('6/12/13', '17:13:19', 69, 64, '18'),
('6/13/13', '18:13:19', 67, 68, '2'),
('6/14/13', '19:13:19', 68, 75, '9'),
('6/15/13', '20:13:19', 69, 58, '9'),
('6/16/13', '21:13:19', 69, 77, '19'),
('6/17/13', '22:13:19', 67, 73, '2'),
```

```
In [21]: %%sql_show
         SElECT buildingID, temp_diff from tempdiffs where temp_diff < 0
```

buildingID	temp_diff
7	-1
17	-5
1	-9
12	-4
8	-3
8	-9
17	-13
17	-4
6	-3
1	-11
17	-4
16	-15
16	-3
5	-12
4	-11
18	-10
4	-7
3	-4
8	-7
6	-10
15	-3
8	-8
16	-3
3	-5
20	-1
8	-10
11	-2

Now finally we need to check the build of this data on the SQL notebook, for that we will just collect the data provided in the first shard.

Total number of records

- We will have have a total of 50 records in the entire file

Command – sqlContect.sql('A,B,C')

```
In [18]: #x = sqlContext.sql(' SELECT buildingID FROM hvac ')

In [19]: x.collect()

Out[19]: [Row(buildingID='2', temp_diff=7, date='6/12/13'),
          Row(buildingID='18', temp_diff=5, date='6/12/13'),
          Row(buildingID='7', temp_diff=-1, date='6/12/13'),
          Row(buildingID='17', temp_diff=-5, date='6/12/13'),
          Row(buildingID='14', temp_diff=7, date='6/12/13'),
          Row(buildingID='20', temp_diff=2, date='6/12/13'),
          Row(buildingID='16', temp_diff=4, date='6/12/13'),
          Row(buildingID='3', temp_diff=5, date='6/12/13'),
          Row(buildingID='1', temp_diff=-9, date='6/12/13'),
          Row(buildingID='17', temp_diff=3, date='6/12/13'),
          Row(buildingID='11', temp_diff=2, date='6/12/13'),
          Row(buildingID='12', temp_diff=-4, date='6/12/13'),
          Row(buildingID='8', temp_diff=-3, date='6/12/13'),
          Row(buildingID='5', temp_diff=13, date='6/12/13'),
          Row(buildingID='8', temp_diff=-9, date='6/12/13'),
          Row(buildingID='2', temp_diff=2, date='6/12/13'),
```

Part II: Experiment with Spark on AWS.

Amazon Elastic Map Reduce (EMR)

➔ EC2(**Compute server**) – This is basically a **regular server** instance that is used to deploy and the required resources over the VM as per the choices made by the users.
- This is used to deploy the VM.
- Manage resources over those VM's.
- Finally to migrate services and monitor volumes.

➔ S3(**Storage utility**) – This is a basic protocol that acts like a storage bucket.
- The prime feature of this protocol is to deal with the data as per service request.
- We need this to store the data in the **data buckets** which are later used to store and move the data across volumes created.

No we need to setup the cluster and run a stream to analyze on top of it.

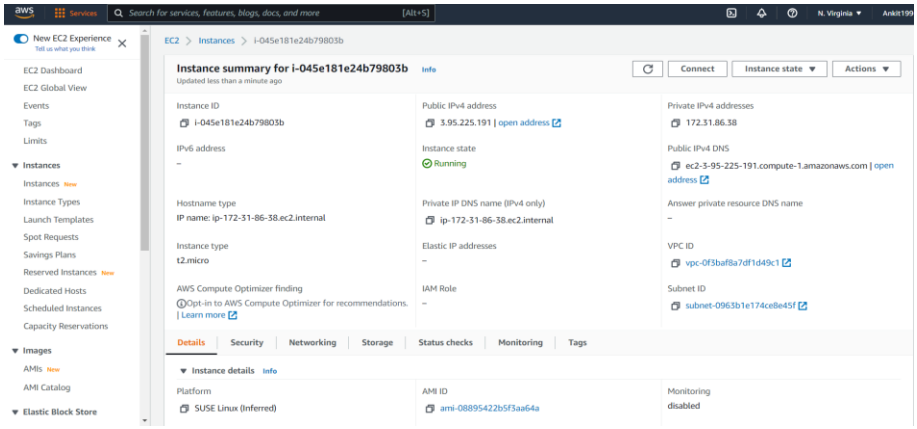

Experimenting with Elastic Map Reduce on the minimal cluster.

Step 2 –

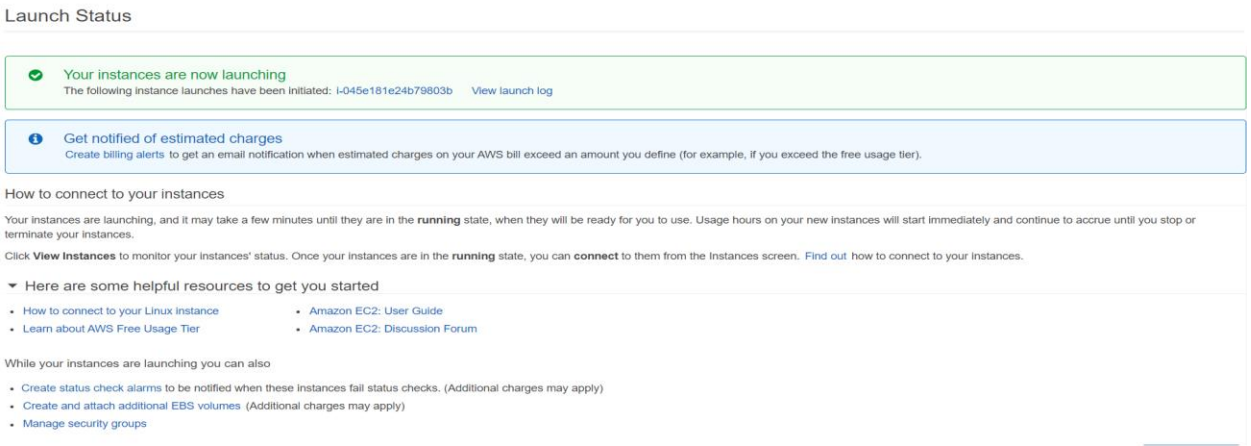

Amazon Cluster Info

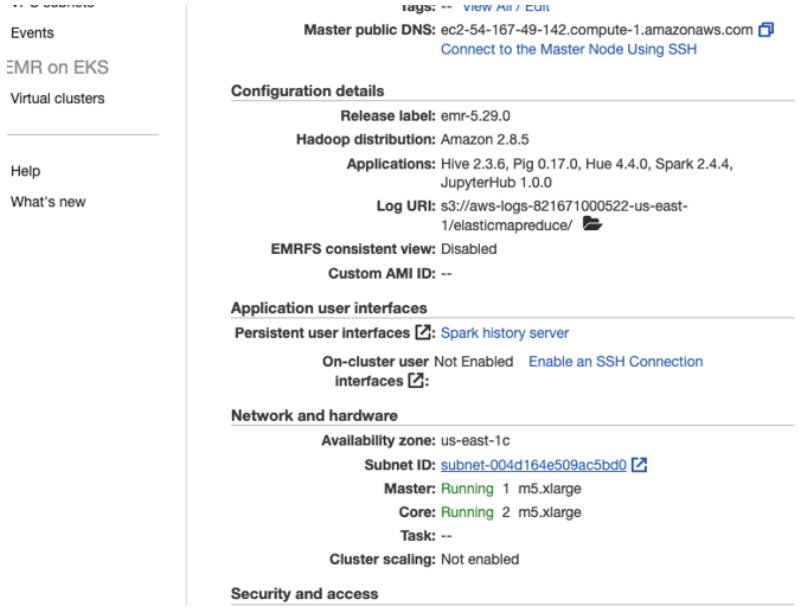

We need to make sure that the stream is ready.

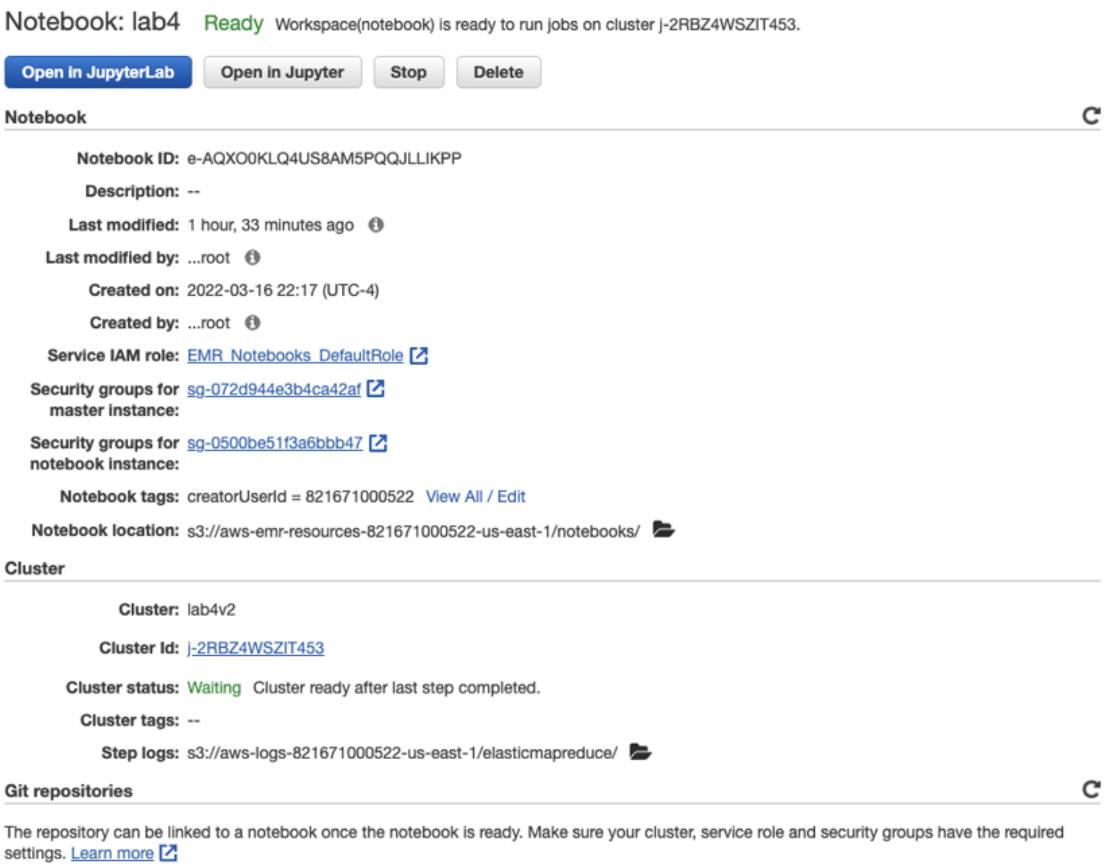

Next we need to begun the EMR and work with the yarn data in livy2.
IP address - **“hdfs://ip-172-31-22-51.ec2.internal:8020/user/wiki/text.txt”**
Text file = **txtfile.repartion(10)**
Array - Finally we need to put these values onto the array in ses2

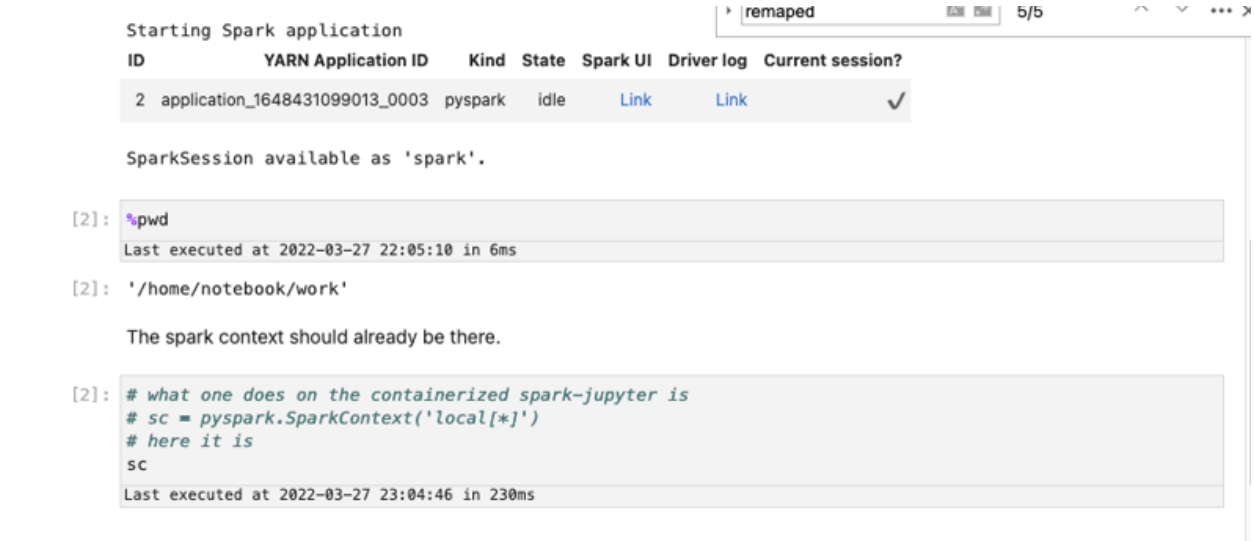

Text file = **txtfile.repartion(10)**
Splitting this test lines into black spaces.

Now finally we will complete the experiment by extracting the heading from each and every line.

[34]:
```
wikidump = sc.textFile("hdfs://ip-172-31-22-51.ec2.internal:8020/user/wiki/text.txt")
```
Last executed at 2022-03-27 23:06:21 in 204ms

[35]:
```
wikidump.count()
```
Last executed at 2022-03-27 23:06:33 in 11.86s

▸ Spark Job Progress

```
21694501
```

[36]:
```
wikidump.getNumPartitions()
```
Last executed at 2022-03-27 23:06:33 in 189ms

```
8
```

[37]:
```
def findtitle(line):
    if line.find('<title>') > -1:
        return True
    else:
        return False
```
Last executed at 2022-03-27 23:06:34 in 285ms

[38]:
```
titles = wikidump.filter(lambda p: findtitle(p))
```
Last executed at 2022-03-27 23:06:34 in 221ms

[39]:
```
titles.count()
```
Last executed at 2022-03-27 23:06:48 in 13.82s

▸ Spark Job Progress

```
0
```

[40]:
```
titles.cache()
```
Last executed at 2022-03-27 23:06:48 in 247ms

Part V: Experiment with Streaming Big Data on AWS; execute and document the following Exercise.

Note – In this case I have used the same json repository mentioned in tutorial as if we take the other one it will shoot up the cost because we have chosen provisioned data. File chosen - .json available with tutorial.

1. Setup the kinesis account on amazon and make a new strem. Stream name – as14128-stream. This is just a dummy stream without any data as of this moment.

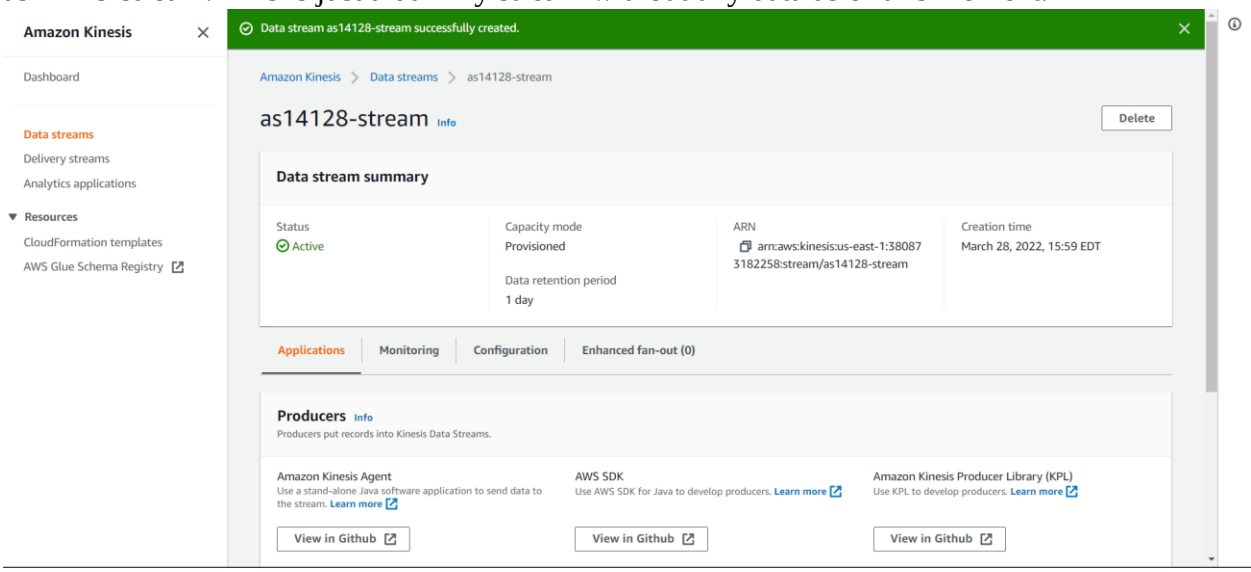

2. Next step is to check if all the things are in place and see that the monitoring **charts** are empty.

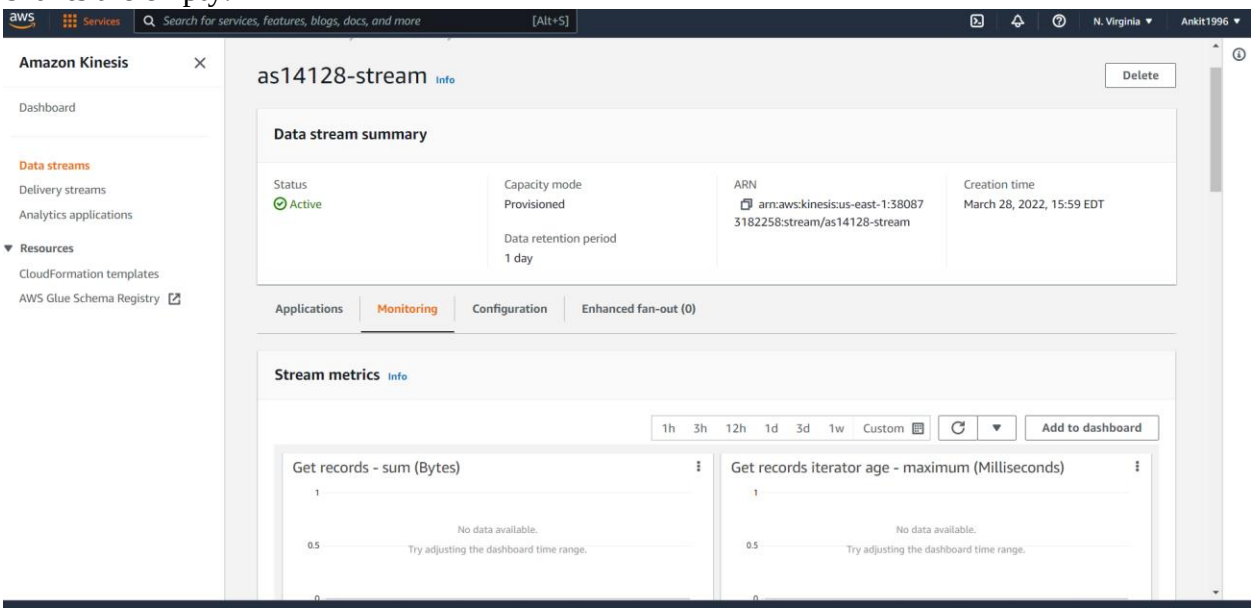

3. Now that we have checked that the charts are empty, we need to go ahead make sure that we have noted don the **configuration** for this stream as it will come in handy later in the process.
   - ➔ Capacity Mode – Provisioned
   - ➔ Total Shards -1 (Only passing a single file onto this stream so no point in getting more)
   - ➔ Write capacity on Stream – 1Mib/Second (The data will come in pretty quick so need to stop to reduce the cost)
   - ➔ Read capacity – 2 Mib/second

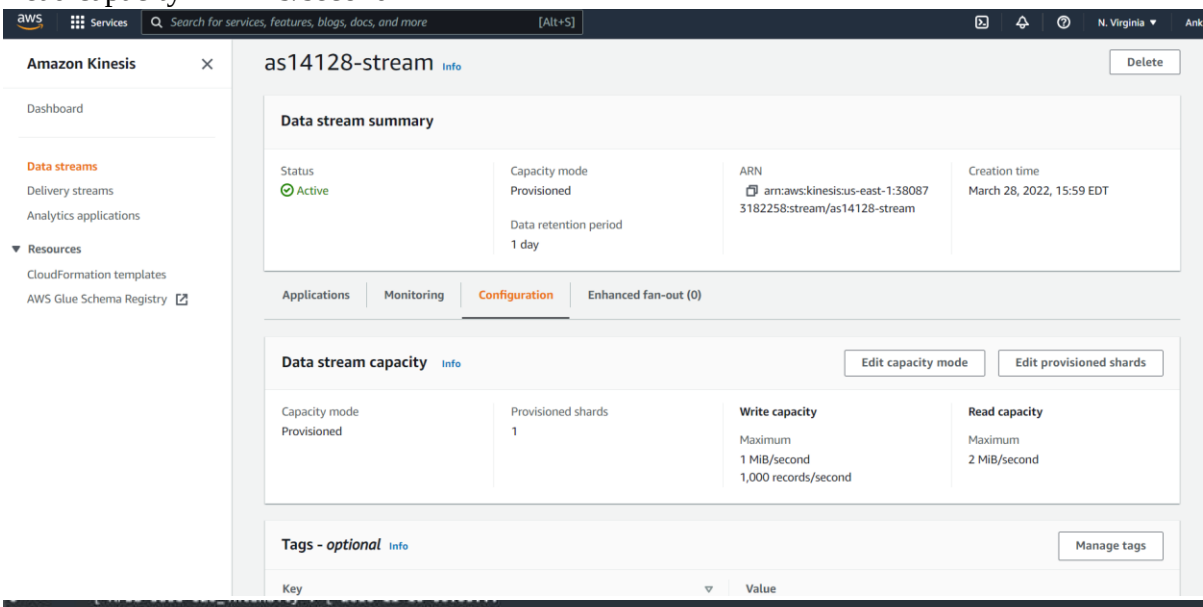

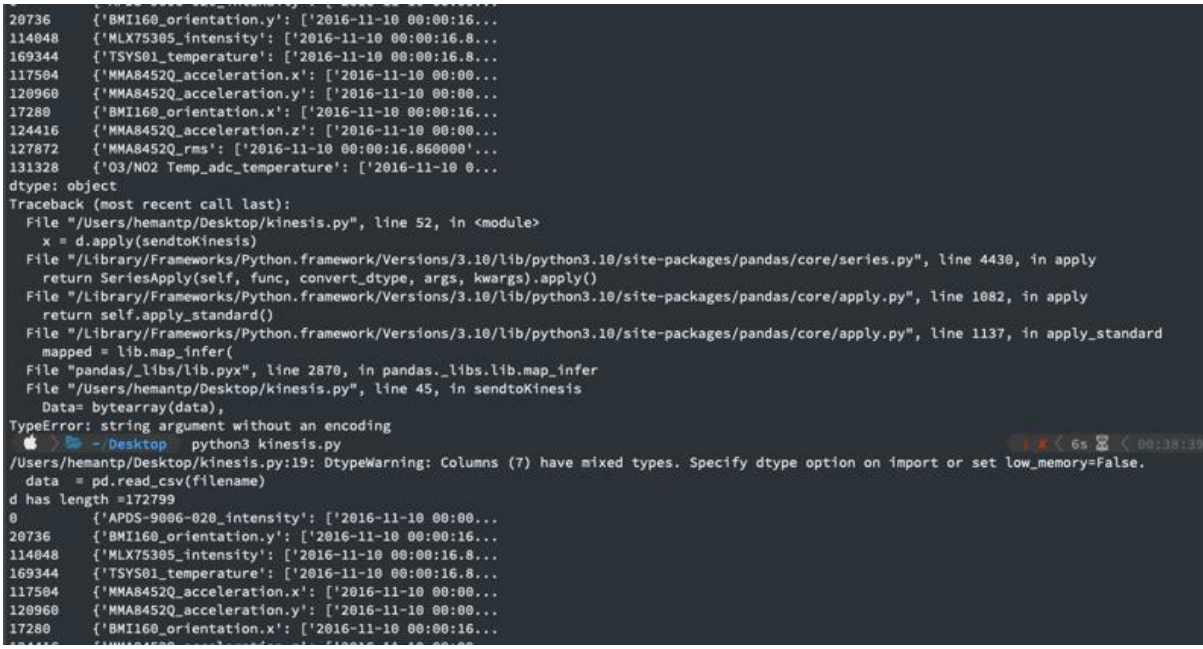

4. Now we need to create a new Stack on top op which we are going to overlay the data. In our case we are taking a small .json file so we will be using the basic stack.
   Creating account on KDG

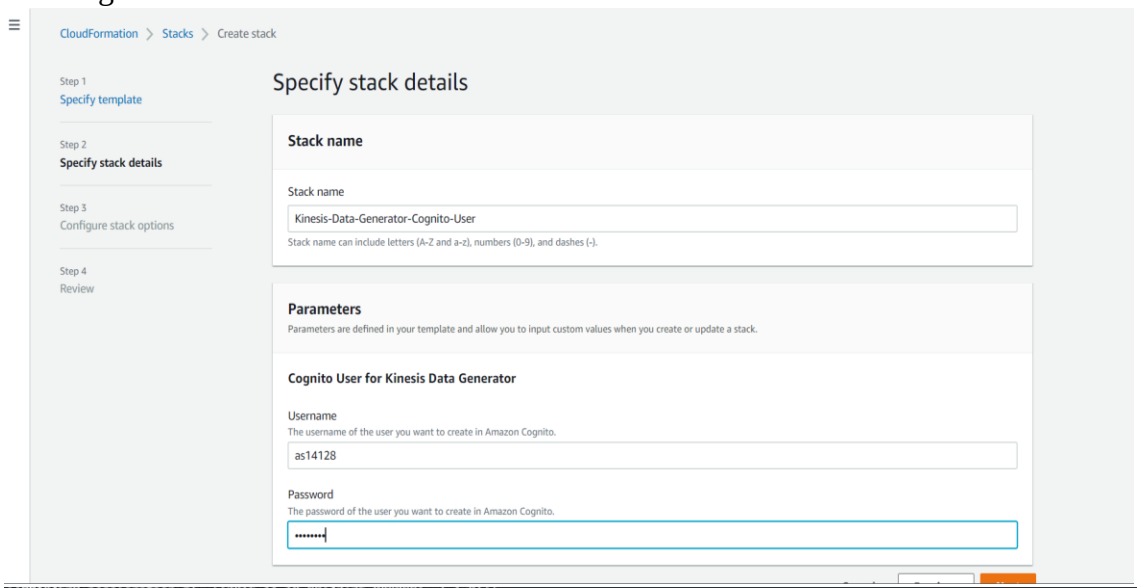

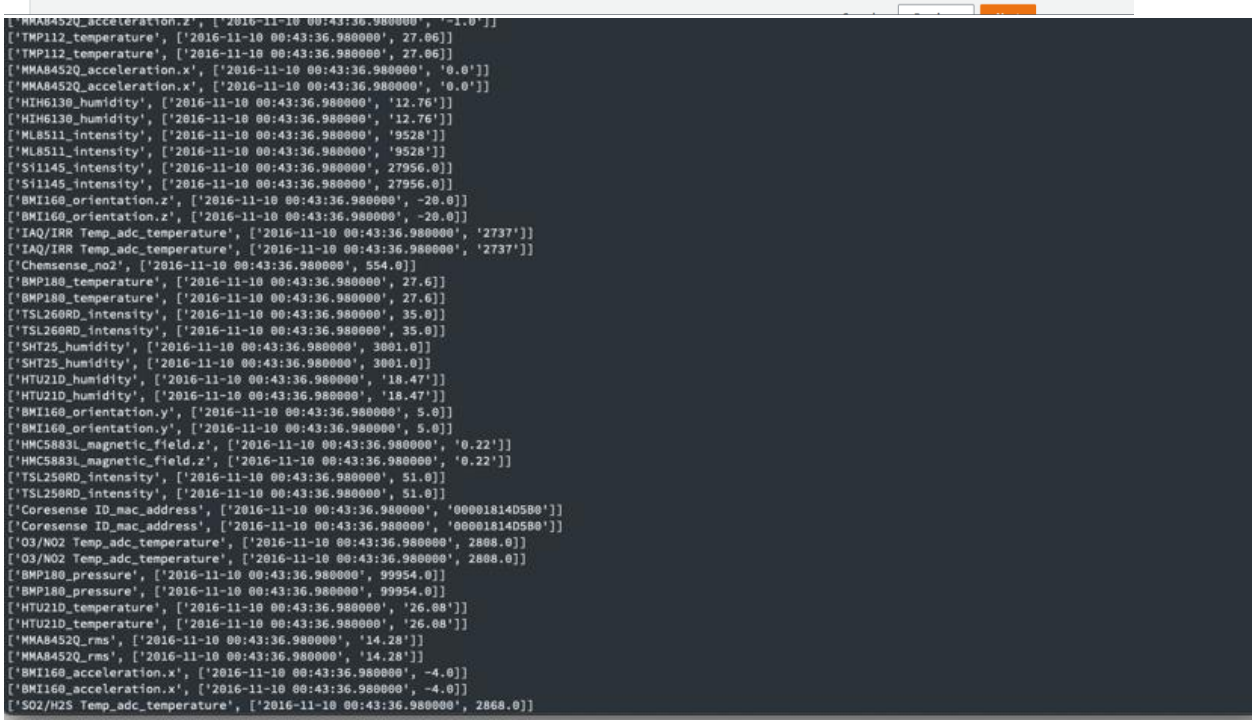

5. Now we need to start creating the stack on this profile and monitor for events to be created.

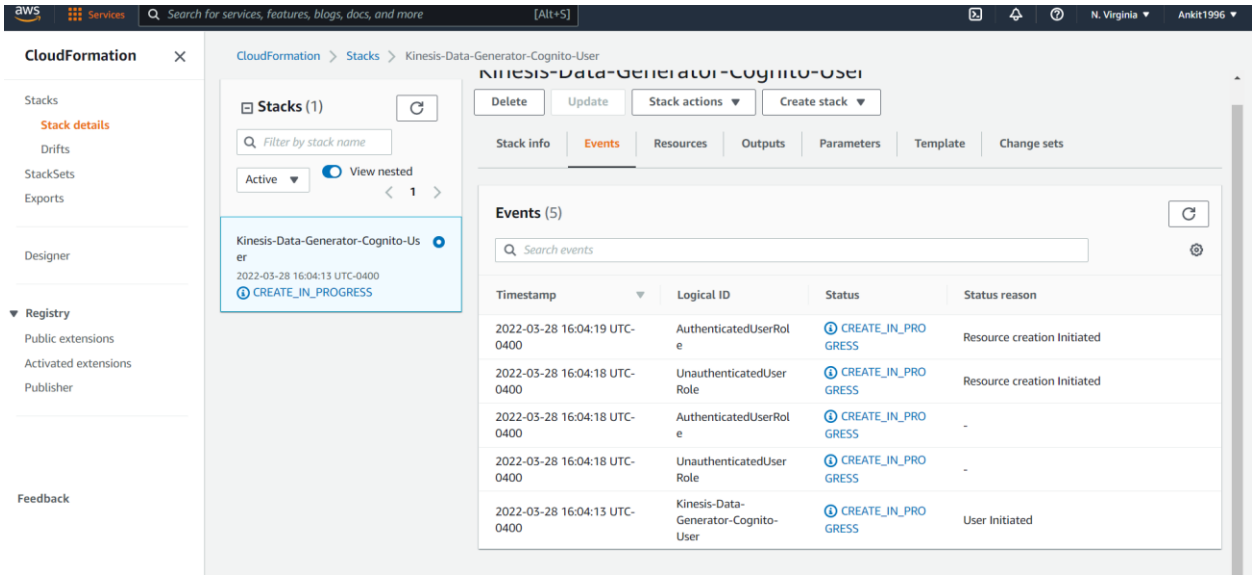

6. Now we need to start passing the .json file onto the stack that we had created on our KDG.
   No we need to be quick because as soon as the data is released the entire burst will come onto the stack and we will need to stop it quick.

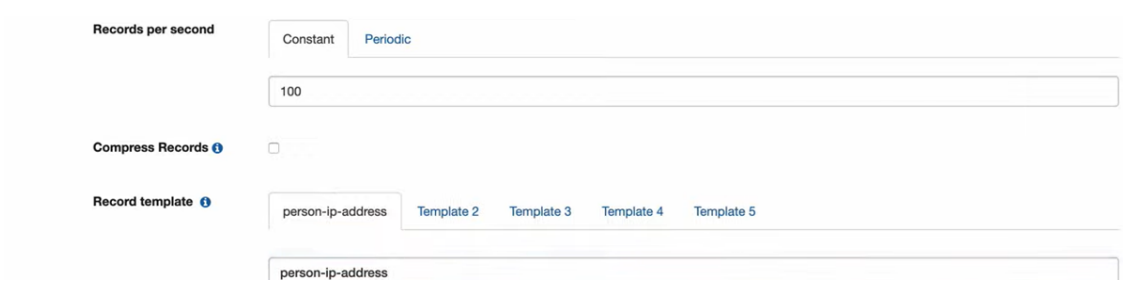

Passing a Valid input depending upon the test file in exercise.

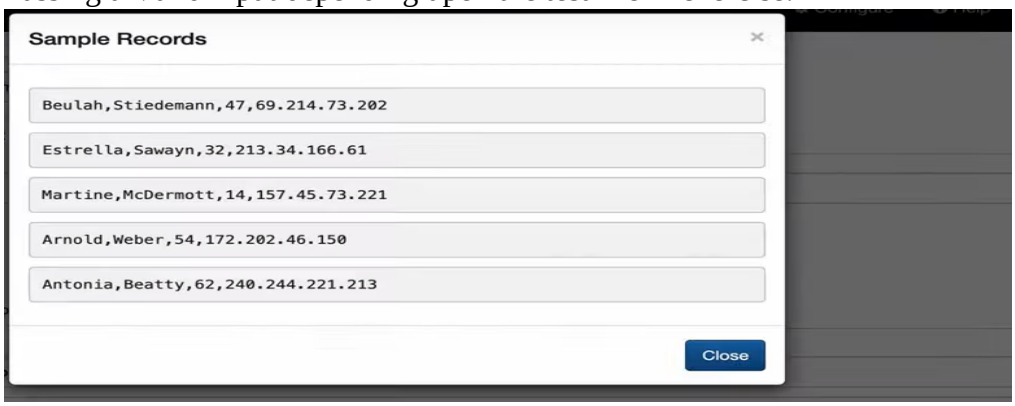

Tracking the progress to monitor the charts later.

7. Now we need to be quick as time is money. The burst has begin and we will soon see the influx of data in the empty charts below.

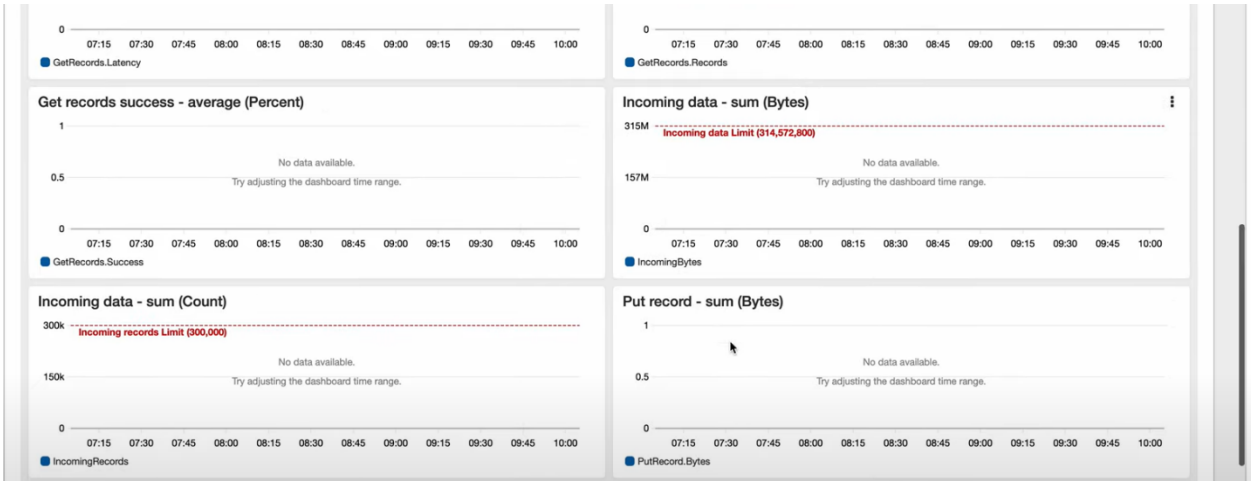

Just to check the influx of data we will need to one of the chart to see if the data overlays properly onto this stack.

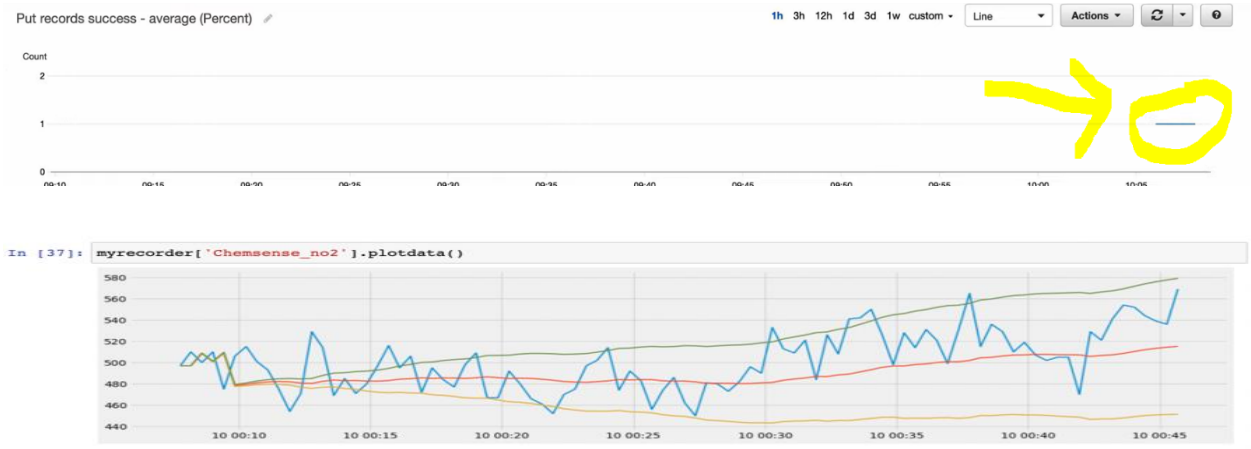

Since this method cost money for as long as it stays up, I have taken it down within a minute post recording the **completion** of the activity.

8. Finally I will be deleting the resources from KDG first.

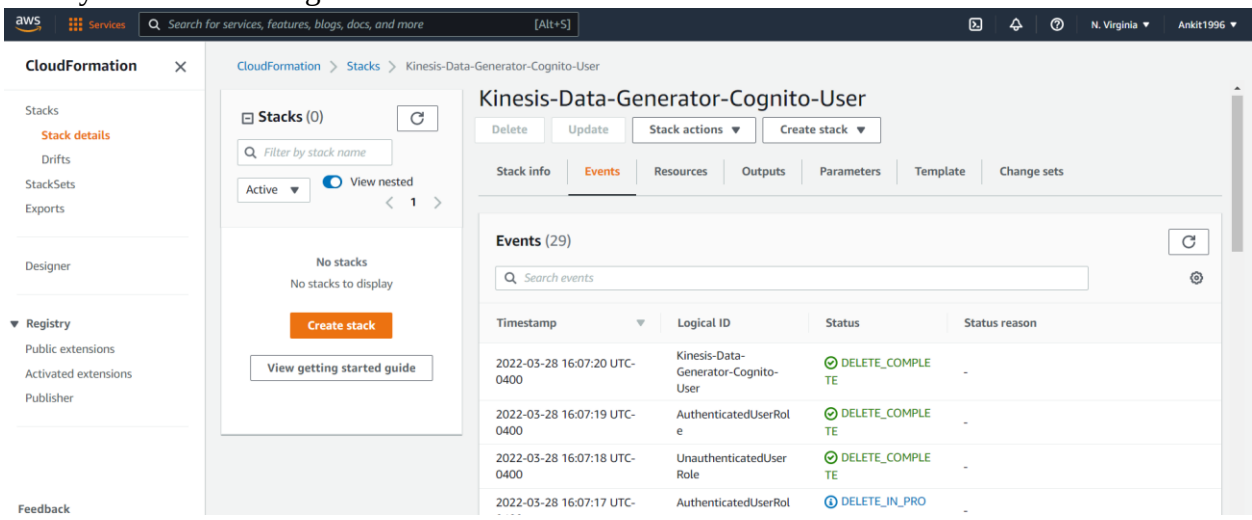

9. Then finally, I will delete the main Kinesis Datastream as well to cut down any further resource utilization.

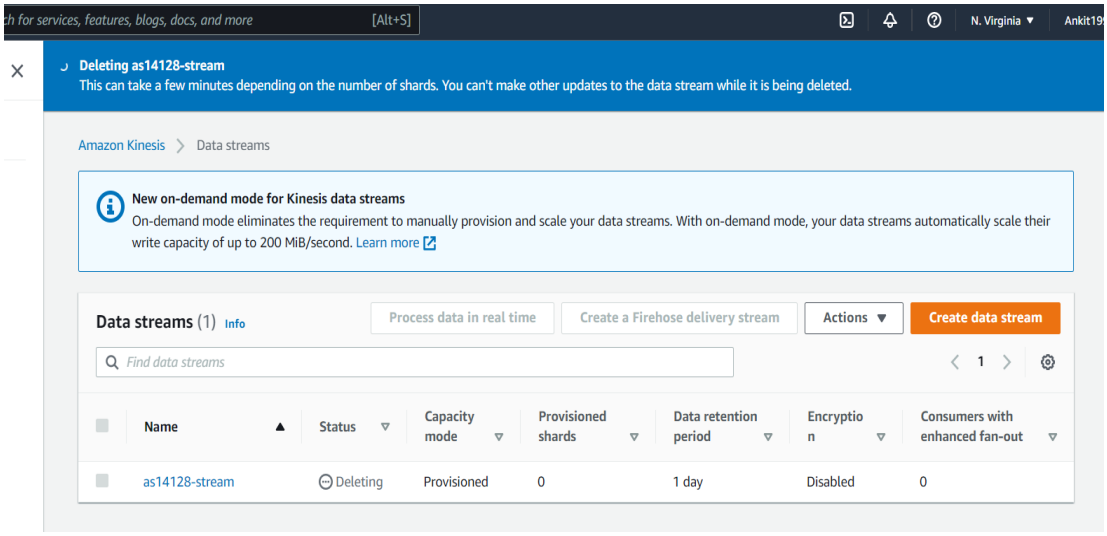

Part III: Google Datalab

1. We need to setup both of the notebooks on the google labs CLI.
   a. Notebook 1

```
ankit@LAPTOP-S2U1QMGB:/mnt/c$ docker run -it -p 8888:8888 dbgannon/tutorial
cp: omitting directory '/tutorial_notebooks/graph'
/home/jovyan/work
total 16
drwxr-xr-x 1 jovyan users 4096 Mar 24 23:31 .
drwxr-xr-x 1 jovyan users 4096 Jul  8  2017 ..
drwxr-xr-x 2 jovyan users 4096 Mar 24 23:31 notebooks
[I 23:31:53.041 NotebookApp] Writing notebook server cookie secret to /home/jovyan/.local/share/jupyter/runtime/notebook_cookie
_secret
[I 23:31:53.714 NotebookApp] JupyterLab alpha preview extension loaded from /opt/conda/lib/python3.5/site-packages/jupyterlab
[I 23:31:53.720 NotebookApp] Serving notebooks from local directory: /home/jovyan/work
[I 23:31:53.720 NotebookApp] 0 active kernels
[I 23:31:53.721 NotebookApp] The Jupyter Notebook is running at: https://[all ip addresses on your system]:8888/
```

   b. Checking both the notebooks in the system GUI.

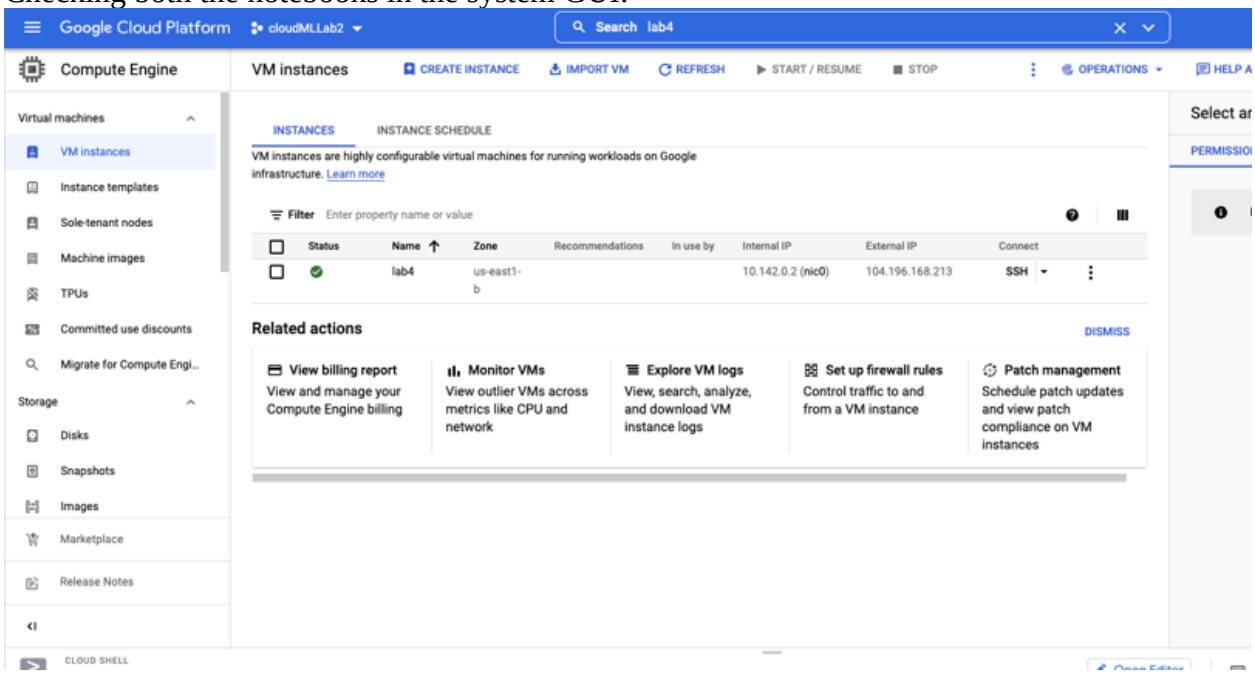

   c. Check the basic lab information required for this experiment.

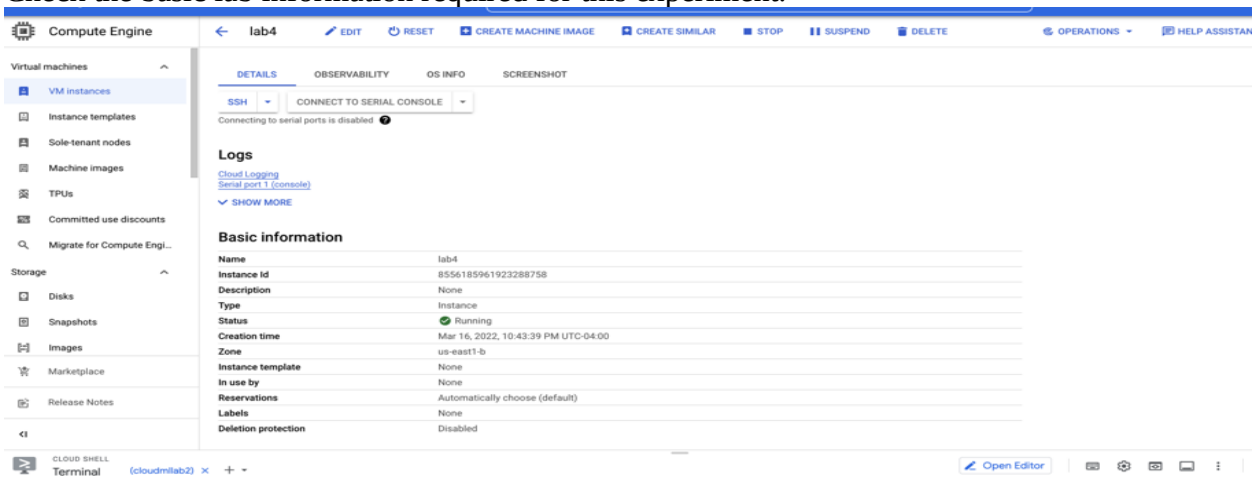

d. Finally uploading both the datalab notebooks to the google cloud.

- **Datalab I**

  In this first lab we are trying to track the rate of spread between the two cities that are separated over a distance of 2000 miles.
  Having said that we need to run this on a single notepook.

```
In [ Saved to this PC lotlib
     import numpy as np
     import matplotlib.pyplot as plot
     %matplotlib inline

In [2]: import datalab.bigquery as bq
```

Checking the same inputs in the reports before commit.

```
In [190]: %bigquery sample --table lookerdata:cdc.project_tycho_reports --count 5
```

Out[190]:

epi_week	state	loc	loc_type	disease	cases	incidence_per_100000
199329	IA	IOWA	STATE	PERTUSSIS	0	0.0
199330	IA	IOWA	STATE	PERTUSSIS	0	0.0
199336	IA	IOWA	STATE	PERTUSSIS	0	0.0
199342	IA	IOWA	STATE	PERTUSSIS	0	0.0
199343	IA	IOWA	STATE	PERTUSSIS	0	0.0

(rows: 5, time: 0.5s, 38MB processed, job: job_DfIrEHhcwSOYXruJNDCgVc4e0cI)

Now we need to track the activity on the graph by importing the "**rcParams**" from the same pyLab.

```
In [287]: from pylab import rcParams
          rcParams['figure.figsize'] = 20, 5
          with plot.style.context('fivethirteight'):
            plot.plot(rwI, rvI, linewidth=2)
            plot.plot(rwW, rvW, linewidth=2)
```

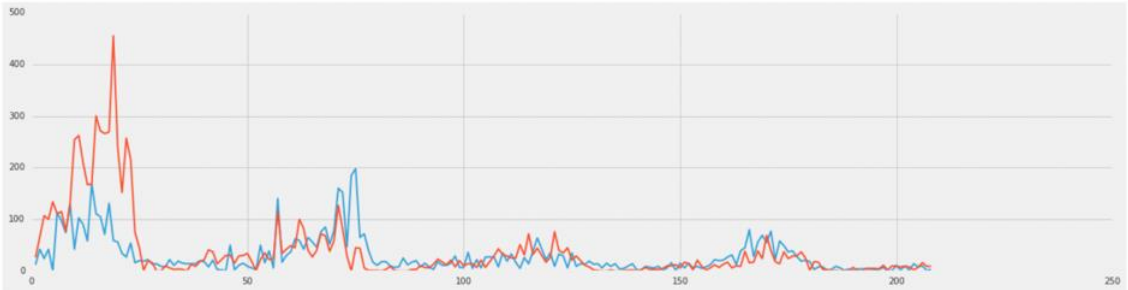

```
In [271]: from pylab import rcParams
          rcParams['figure.figsize'] = 20, 5
          with plot.style.context('fivethirtyeight'):
            plot.plot(realweekI, rubelINval, linewidth=2)
            plot.plot(realweekW, rubelWAval, linewidth=2)
```

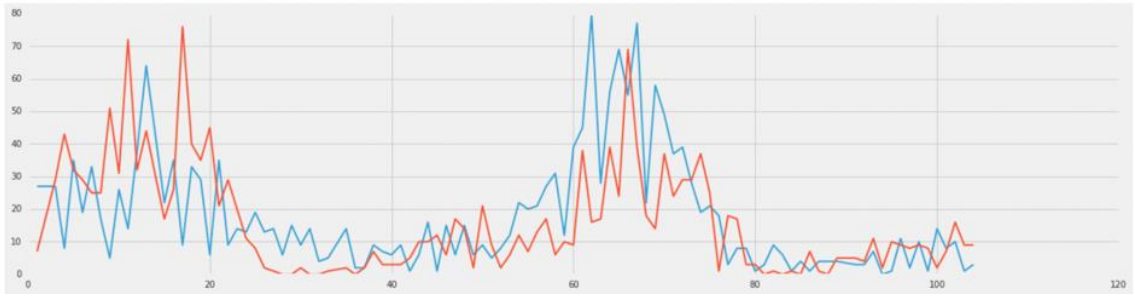

Next we need to compare the two years. The data stored in lab 1 , will go through the same normalization and plot the boundaries for the next 4 years.

- **Datalab II** – Abnormality in the weather stations in Washington.

```
In [140]: %%sql
          SELECT
            max, (max-32)*5/9 celsius, mo, da, state, stn, name
          FROM (
            SELECT
              max, mo, da, state, stn, name
            FROM
              [bigquery-public-data:noaa_gsod.gsod2015] a
            JOIN
              [bigquery-public-data:noaa_gsod.stations] b
            ON
              a.stn=b.usaf
              AND a.wban=b.wban
            WHERE
              state="WA"
              AND max<1000
              AND country='US' )
          ORDER BY
            max DESC
```

Out[140]:

max	celsius	mo	da	state	stn	name
113.0	45.0	06	29	WA	727846	WALLA WALLA RGNL
113.0	45.0	06	28	WA	727846	WALLA WALLA RGNL

By tracing the below graph, we can see that the skagit station is the culprit.
Apart from this we can see the anomalies in the month of September and March for the same year.

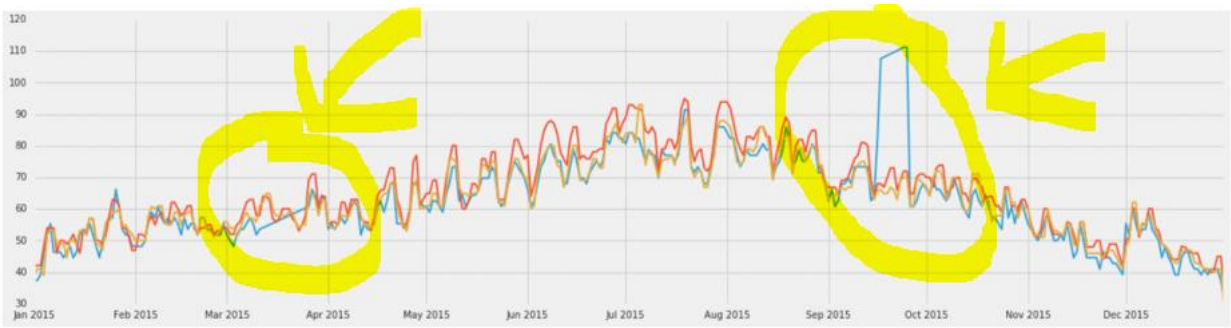

Part VI: Google Datalab

This experiment was done on an EMR Hadoop cluster as well. A program was executed as a spark batch processing job to count the sum of the first 1 million prime numbers

Python code used is mentioned below.

```
if __name__ == "__main__":
    conf = SparkConf().setAppName("primeNumbers").setMaster("local[*]")
    sc = SparkContext(conf = conf)

    lines = sc.textFile("1m.csv")
    header = lines.first() #extract header
    lines = lines.filter(lambda row:  row != header)

    intNumbers = lines.map(lambda line: int(line.split(",")[1]))

    print(intNumbers.take(10))

    print("Sum is: {}".format(intNumbers.reduce(lambda x, y: x + y)))
[hadoop@ip-172-31-22-51 ~]$
```

Final build and answer of the above .py file.

```
22/03/28 05:57:37 INFO TaskSchedulerImpl: Removed TaskSet 2.0, whose tasks have all completed, from pool
22/03/28 05:57:37 INFO DAGScheduler: ResultStage 2 (reduce at /home/hadoop/sumOfPrimes.py:15) finished in 2.663 s
22/03/28 05:57:37 INFO DAGScheduler: Job 2 finished: reduce at /home/hadoop/sumOfPrimes.py:15, took 2.666297 s
Sum is: 7472951481636
22/03/28 05:57:37 INFO SparkContext: Invoking stop() from shutdown hook
22/03/28 05:57:37 INFO SparkUI: Stopped Spark web UI at http://ip-172-31-22-51.ec2.internal:4040
22/03/28 05:57:37 INFO MapOutputTrackerMasterEndpoint: MapOutputTrackerMasterEndpoint stopped!
22/03/28 05:57:37 INFO MemoryStore: MemoryStore cleared
22/03/28 05:57:37 INFO BlockManager: BlockManager stopped
22/03/28 05:57:37 INFO BlockManagerMaster: BlockManagerMaster stopped
22/03/28 05:57:37 INFO OutputCommitCoordinator$OutputCommitCoordinatorEndpoint: OutputCommitCoordinator stopped!
22/03/28 05:57:37 INFO SparkContext: Successfully stopped SparkContext
22/03/28 05:57:37 INFO ShutdownHookManager: Shutdown hook called
22/03/28 05:57:37 INFO ShutdownHookManager: Deleting directory /mnt/tmp/spark-9925eeda-90f4-447c-a889-3a45a91d2c03
22/03/28 05:57:37 INFO ShutdownHookManager: Deleting directory /mnt/tmp/spark-d1d76190-52b0-4ec8-9907-aef0abd0ea3a/pyspark-150fb1c5-50bd-4a51-bc65-
fa80beb0e57b
22/03/28 05:57:37 INFO ShutdownHookManager: Deleting directory /mnt/tmp/spark-d1d76190-52b0-4ec8-9907-aef0abd0ea3a
```

Part VII: Build additional experiments to document the use of MapReduce using Hadoop.

1. Create a generic mapper function that calculates the number of words in the text file put in hdfs.

```
#!/usr/bin/env python3
import sys
import string
for line in sys.stdin:
    line = line.strip()
    words = line.split()
    for w in words:
        table = w.maketrans('', '', string.punctuation)
        w = w.translate(table).lower()
        print(w, '\t', 1)
~
~
~
~
```

2. Now we need to put the reducer file to count the number of words in the document.

```
#!/usr/bin/env python3
from collections import defaultdict
import sys
word_count = defaultdict(int)
for line in sys.stdin:
    try:
        line = line.strip()
        word, count = line.split()
        count = int(count)
    except:
        continue
word_count[word] += count
for word, count in word_count.items():
    print(word, count)
~
```

3. Submit both the application and monitor the status of map reduce over the files.

```
22/03/28 04:33:55 INFO mapreduce.Job: Running job: job_1648431099013_0004
22/03/28 04:34:01 INFO mapreduce.Job: Job job_1648431099013_0004 running in uber mode : false
22/03/28 04:34:01 INFO mapreduce.Job:  map 0% reduce 0%
22/03/28 04:34:07 INFO mapreduce.Job:  map 2% reduce 0%
22/03/28 04:34:08 INFO mapreduce.Job:  map 4% reduce 0%
22/03/28 04:34:12 INFO mapreduce.Job:  map 6% reduce 0%
22/03/28 04:34:13 INFO mapreduce.Job:  map 16% reduce 0%
22/03/28 04:34:17 INFO mapreduce.Job:  map 20% reduce 0%
22/03/28 04:34:21 INFO mapreduce.Job:  map 22% reduce 0%
22/03/28 04:34:22 INFO mapreduce.Job:  map 33% reduce 0%
22/03/28 04:34:25 INFO mapreduce.Job:  map 35% reduce 0%
22/03/28 04:34:27 INFO mapreduce.Job:  map 37% reduce 0%
22/03/28 04:34:29 INFO mapreduce.Job:  map 39% reduce 0%
22/03/28 04:34:31 INFO mapreduce.Job:  map 47% reduce 0%
22/03/28 04:34:38 INFO mapreduce.Job:  map 49% reduce 0%
22/03/28 04:34:39 INFO mapreduce.Job:  map 53% reduce 0%
22/03/28 04:34:40 INFO mapreduce.Job:  map 55% reduce 0%
22/03/28 04:34:45 INFO mapreduce.Job:  map 57% reduce 0%
22/03/28 04:34:46 INFO mapreduce.Job:  map 57% reduce 6%
22/03/28 04:34:48 INFO mapreduce.Job:  map 59% reduce 6%
22/03/28 04:34:49 INFO mapreduce.Job:  map 63% reduce 6%
22/03/28 04:34:51 INFO mapreduce.Job:  map 65% reduce 6%
```

4. Checking the count of a single sentence.

```
1351.txt  1360.txt  1370.txt  1378.txt  1389.txt  1397.txt
[hadoop@ip-172-31-22-51 ~]$ ls -lh books-input
total 0
[hadoop@ip-172-31-22-51 ~]$ cd books-input/
[hadoop@ip-172-31-22-51 books-input]$ mv ../1*
mv: target '../1398.txt' is not a directory
[hadoop@ip-172-31-22-51 books-input]$ mv ../1* .
[hadoop@ip-172-31-22-51 books-input]$ ls
1340.txt  1347.txt  1354.txt  1360.txt  1366.txt  1372.txt  1377.txt  1385.txt  1390.txt  1395.txt
1341.txt  1348.txt  1356.txt  1361.txt  1367.txt  1373.txt  1378.txt  1386.txt  1391.txt  1396.txt
1343.txt  1351.txt  1357.txt  1362.txt  1369.txt  1374.txt  1379.txt  1387.txt  1392.txt  1397.txt
1344.txt  1352.txt  1358.txt  1363.txt  1370.txt  1375.txt  1380.txt  1388.txt  1393.txt  1398.txt
1345.txt  1353.txt  1359.txt  1364.txt  1371.txt  1376.txt  1384.txt  1389.txt  1394.txt
[hadoop@ip-172-31-22-51 books-input]$ cd ..
[hadoop@ip-172-31-22-51 ~]$ vim mapper.py
[hadoop@ip-172-31-22-51 ~]$ chmod +x mapper.py
[hadoop@ip-172-31-22-51 ~]$ printf 'My name is Karim\nWhat is your name' | ./mapper.py
my      1
name    1
is      1
karim   1
what    1
is      1
your    1
name    1
```

5. Matching results of the book. (Initial file)

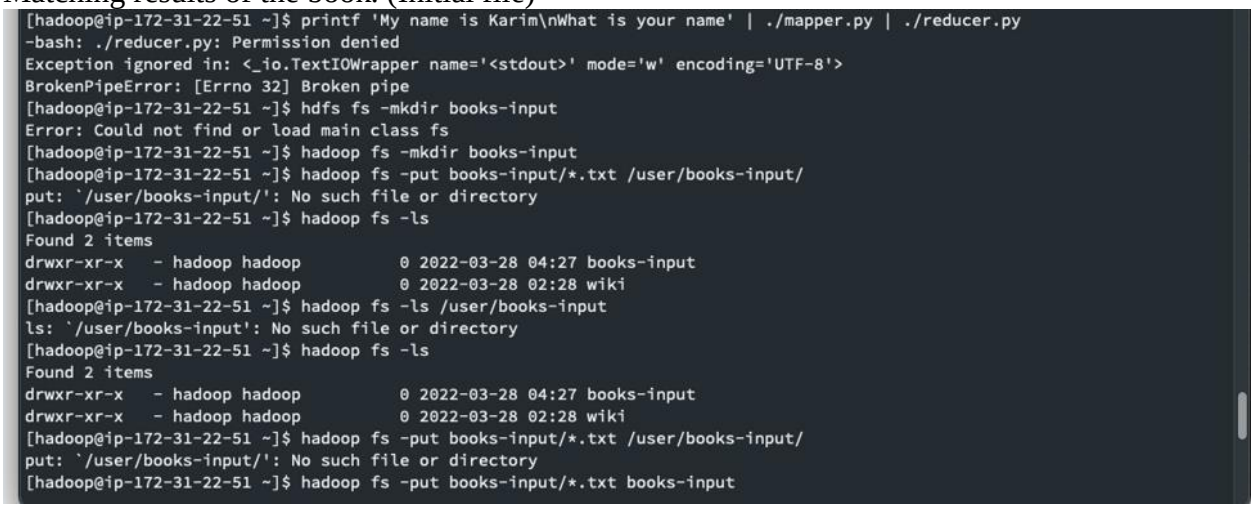

```
[hadoop@ip-172-31-22-51 ~]$ printf 'My name is Karim\nWhat is your name' | ./mapper.py | ./reducer.py
-bash: ./reducer.py: Permission denied
Exception ignored in: <_io.TextIOWrapper name='<stdout>' mode='w' encoding='UTF-8'>
BrokenPipeError: [Errno 32] Broken pipe
[hadoop@ip-172-31-22-51 ~]$ hdfs fs -mkdir books-input
Error: Could not find or load main class fs
[hadoop@ip-172-31-22-51 ~]$ hadoop fs -mkdir books-input
[hadoop@ip-172-31-22-51 ~]$ hadoop fs -put books-input/*.txt /user/books-input/
put: `/user/books-input/': No such file or directory
[hadoop@ip-172-31-22-51 ~]$ hadoop fs -ls
Found 2 items
drwxr-xr-x   - hadoop hadoop          0 2022-03-28 04:27 books-input
drwxr-xr-x   - hadoop hadoop          0 2022-03-28 02:28 wiki
[hadoop@ip-172-31-22-51 ~]$ hadoop fs -ls /user/books-input
ls: `/user/books-input': No such file or directory
[hadoop@ip-172-31-22-51 ~]$ hadoop fs -ls
Found 2 items
drwxr-xr-x   - hadoop hadoop          0 2022-03-28 04:27 books-input
drwxr-xr-x   - hadoop hadoop          0 2022-03-28 02:28 wiki
[hadoop@ip-172-31-22-51 ~]$ hadoop fs -put books-input/*.txt /user/books-input/
put: `/user/books-input/': No such file or directory
[hadoop@ip-172-31-22-51 ~]$ hadoop fs -put books-input/*.txt books-input
```